

# Exploratory Data Analysis of Haberman's Dataset

## Dataset Information:

The dataset contains cases from a study that was conducted between 1958 and 1970 at the University of Chicago's Billings Hospital on the survival of patients who had undergone surgery for breast cancer.

## Attribute Information:

1. Age of patient at time of operation (numerical)
2. Patient's year of operation (year - 1900, numerical)
3. Number of positive axillary nodes detected (numerical)
4. Survival status (1 - The patient survived 5 years or longer, 2 - The patient died within 5 years)

## Objective:

The objective here is to classify whether or not a patient will survive for 5 years following a surgery for breast cancer based on the age of the patient, the year of operation and the number of positive axillary nodes detected.

## Importing and Cleaning of Data:

In this section, the data and the required libraries are imported and made ready for analysis.

In [1]:

```
# Importing the required library stack with their usual aliases

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

import warnings # In order to suppress any warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
# Saved the csv file locally.

hb = pd.read_csv('haberman.csv')

hb.head()
```

Out[2]:

|   | Age | Operation_Year | Axillary_Nodes | Survival_Status |
|---|-----|----------------|----------------|-----------------|
| 0 | 30  | 64             | 1              | 1               |
| 1 | 30  | 62             | 3              | 1               |
| 2 | 30  | 65             | 0              | 1               |
| 3 | 31  | 59             | 2              | 1               |
| 4 | 31  | 65             | 4              | 1               |

In [3]:

```
# Usually, it is good practice to check for any null values in the data before proceeding with the analysis.
Empty_Check = hb[hb.isnull().any(axis = 1)]
Empty_Check.empty
```

Out[3]:

True

In [4]:

```
# As we can see in the dataset, the survival statuses of the patients are represented numerically.
# Let us assign a better notation.
# A 'Yes' instead of a 1 signifying that the patient survived for over 5 years after the surgery.
# A 'No' instead of a 2 signifying that the patient survived for lesser than 5 years after the surgery.

hb['Survival_Status'] = hb['Survival_Status'].map({1:'Yes',2:'No'})
hb.head(10)
```

Out[4]:

|   | Age | Operation_Year | Axillary_Nodes | Survival_Status |
|---|-----|----------------|----------------|-----------------|
| 0 | 30  | 64             | 1              | Yes             |
| 1 | 30  | 62             | 3              | Yes             |
| 2 | 30  | 65             | 0              | Yes             |
| 3 | 31  | 59             | 2              | Yes             |
| 4 | 31  | 65             | 4              | Yes             |
| 5 | 33  | 58             | 10             | Yes             |
| 6 | 33  | 60             | 0              | Yes             |
| 7 | 34  | 59             | 0              | No              |
| 8 | 34  | 66             | 9              | No              |
| 9 | 34  | 58             | 30             | Yes             |

## High Level Statistics:

We answer some really basic questions about the data here like number of features, number of classes, number of datapoints etc.

In [5]:

```
# Shape of the data
print(hb.shape)

print('\n')

# Number of classes
print(hb['Survival_Status'].unique())

print('\n')

# Number of data points per class
print(hb['Survival_Status'].value_counts())

print('\n')

# General description of the data
print(hb.describe().round(2))
```

(306, 4)

['Yes' 'No']

```
Yes      225
No        81
Name: Survival_Status, dtype: int64
```

Age Operation\_Year Axillary\_Nodes

|       | Age    | Operation_Year | Survival_Status |
|-------|--------|----------------|-----------------|
| count | 306.00 | 306.00         | 306.00          |
| mean  | 52.46  | 62.85          | 4.03            |
| std   | 10.80  | 3.25           | 7.19            |
| min   | 30.00  | 58.00          | 0.00            |
| 25%   | 44.00  | 60.00          | 0.00            |
| 50%   | 52.00  | 63.00          | 1.00            |
| 75%   | 60.75  | 65.75          | 4.00            |
| max   | 83.00  | 69.00          | 52.00           |


## Univariate Analysis:

Here, we plot PDFs, CDFs, Box Plots on each feature and try to come to conclusions about the data and see which feature is useful in classification.

In [6]:


```
# 1.
# Plotting PDFs
# PDFs are basically a smoothened version of a histogram.
# They represent the probability that the variable takes a particular value.


sns.FacetGrid(hb, hue='Survival_Status', size = 5) \
    .map(sns.distplot, 'Age') \
    .add_legend()
plt.title('Variation of Survival Status with Age')
plt.show()
```



In [7]:


```
sns.FacetGrid(hb, hue='Survival_Status', size = 5) \
    .map(sns.distplot, 'Operation_Year') \
    .add_legend()
plt.title('Variation of Survival Status with Operation Year')
plt.show()
```





In [8]:

```
sns.FacetGrid(hb, hue='Survival_Status',size = 5) \
    .map(sns.distplot,'Axillary_Nodes') \
    .add_legend()
plt.title('Variation of Survival Status with Axillary Nodes')
plt.show()
```



In [9]:

```
# In order to understand the data in more detail and how each feature
# is impacting the classification, we create a subset of the main dataframe.
# One subset has all the datapoints that lead to a 'Yes' and the other subset
# has all the datapoints that lead to a 'No'.

SS_Yes = hb[hb['Survival_Status'] == 'Yes']
SS_No = hb[hb['Survival_Status'] == 'No']
```

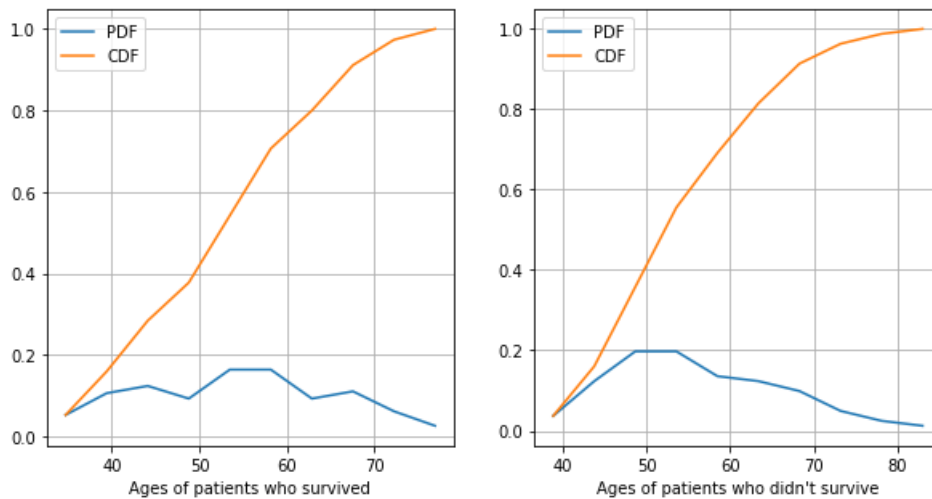
In [10]:

```
# 2.
# Plotting CDFs
# CDFs allow us to determine what percent of the datapoints in a dataset are below a particular value.
plt.figure(figsize = (10,5))

plt.subplot(1,2,1)
counts , bin_edges = np.histogram(SS_Yes['Age'], density = True)
pdf = counts/(sum(counts))
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,label='PDF')
plt.plot(bin_edges[1:], cdf,label='CDF')
plt.grid()
plt.legend()
plt.xlabel('Ages of patients who survived')

plt.subplot(1,2,2)
counts , bin_edges = np.histogram(SS_No['Age'], density = True)
pdf = counts/(sum(counts))
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,label='PDF')
plt.plot(bin_edges[1:], cdf, label='CDF')
plt.grid()
plt.legend()
plt.xlabel('Ages of patients who didn\'t survive')
```

```
plt.show()
```




```
In [11]:
```

```
plt.figure(figsize = (10,5))

plt.subplot(1,2,1)
counts , bin_edges = np.histogram(SS_Yes['Operation_Year'], density = True)
pdf = counts/(sum(counts))
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,label='PDF')
plt.plot(bin_edges[1:], cdf,label='CDF')
plt.grid()
plt.legend()
plt.xlabel('Operation Years of patients who survived')

plt.subplot(1,2,2)
counts , bin_edges = np.histogram(SS_No['Operation_Year'], density = True)
pdf = counts/(sum(counts))
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,label='PDF')
plt.plot(bin_edges[1:], cdf,label='CDF')
plt.grid()
plt.legend()
plt.xlabel('Operation Years of patients who didn\'t survive')

plt.show()
```



```
In [12]:
```



```
plt.figure(figsize = (10,5))

plt.subplot(1,2,1)
counts , bin_edges = np.histogram(SS_Yes['Axillary_Nodes'], density = True)
pdf = counts/(sum(counts))
cdf = np.cumsum(pdf)
```

```
plt.plot(bin_edges[1:],pdf,label='PDF')
plt.plot(bin_edges[1:], cdf,label='CDF')
plt.grid()
plt.legend()
plt.xlabel('Axillary Nodes in patients who survived')

plt.subplot(1,2,2)
counts , bin_edges = np.histogram(SS_No['Axillary_Nodes'], density = True)
pdf = counts/(sum(counts))
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,label='PDF')
plt.plot(bin_edges[1:], cdf,label='CDF')
plt.grid()
plt.legend()
plt.xlabel('Axillary Nodes in patients who didn\'t survive')

plt.show()
```






In [13]:

```
# 3.
# BOXPLOTS are a more visual and intuitive way of visualizing the dataset.
# There are 5 main lines in a box plot.
# The 3 lines in the middle represent the Inter-Quartile_Range, i.e. the 25th, 50th and 75th percentile respectively.
# The lines at the bottom and the top are the whiskers which are the [Q1 - 1.5(IQR)] and the [Q3 + 1.5(IQR)] respectively.
plt.figure(figsize = (15,5))

plt.subplot(131)
sns.boxplot(x='Survival_Status', y = 'Age', data = hb)
plt.subplot(132)
sns.boxplot(x='Survival_Status', y = 'Operation_Year', data = hb)
plt.subplot(133)
sns.boxplot(x='Survival_Status', y = 'Axillary_Nodes', data = hb)

plt.show()
```

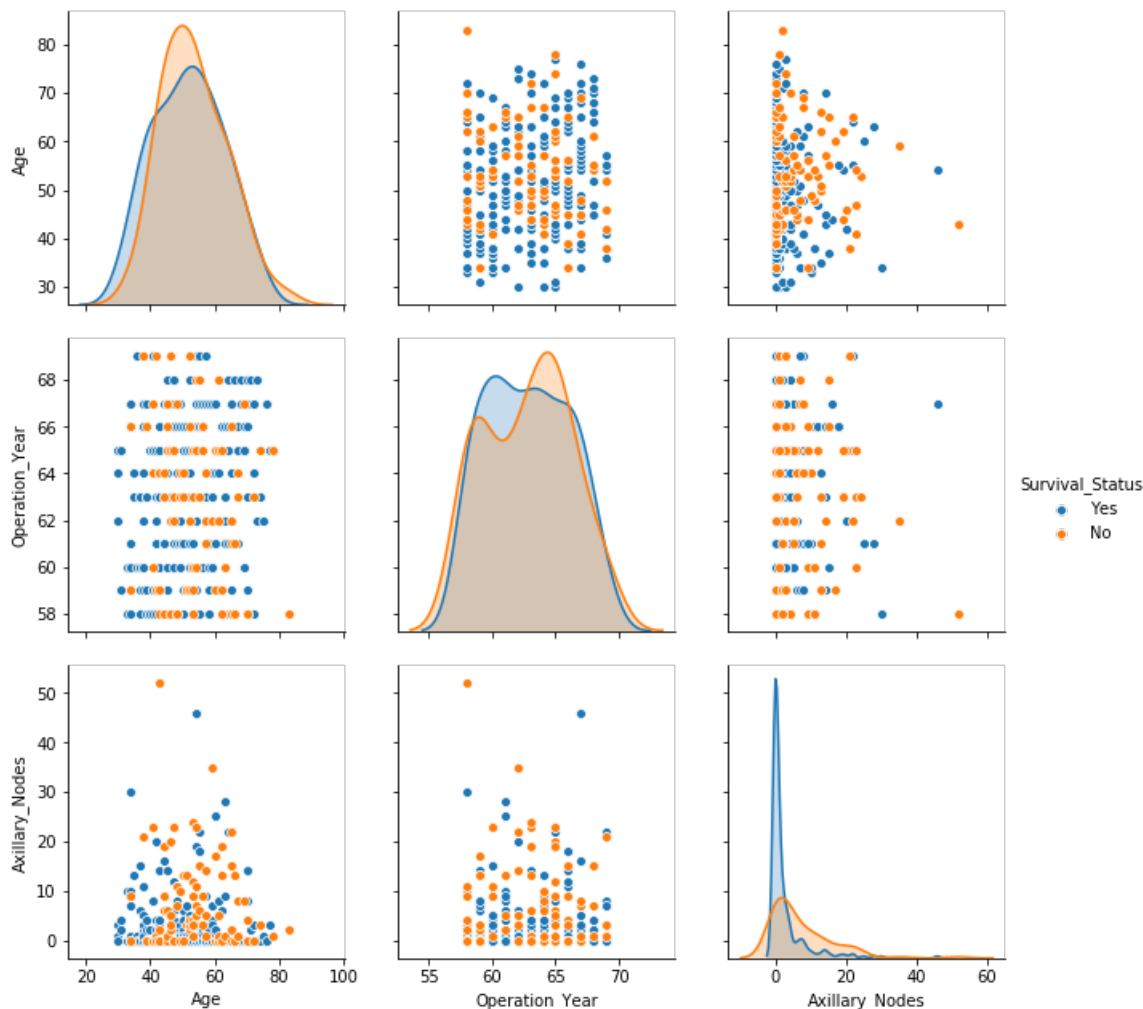


## Bivariate Analysis

Here we plot a pairplot to see which combination of features does a better job at classification

In [14]:

```
# Pairplot
sns.pairplot(hb,hue='Survival_Status',size = 3)
plt.show()
```



## Final Conclusions

1. The number of rows (records of patients) are 306.
2. The number of classes are two, namely 'Yes' and 'No'.
3. There are 225 datapoints that result in a 'Yes' and 81 datapoints that result in a 'No'.
4. This is an imbalanced dataset.
5. The majority of patients that survived the treatment are in the ages between 50-60.
6. The majority of patients that didn't survive the treatment are around 50 years.
7. The number of Axillary Nodes found in patients are usually around (0-5) irrespective of their Survival Status.
8. About 90% of the patients that survived the surgery had less than 10 Axillary Nodes.
9. Roughly 70% of the patients that didn't survive the surgery had less than 10 Axillary Nodes.
10. In the case of Axillary Nodes, there seem to be quite a few number of outliers.
11. The ages of 50% of the patients who survived the treatment lie in between 43-60 roughly.
12. The ages of 50% of the patients who didn't survive the treatment are in between 46- 61 roughly.
13. Although there isn't much difference, the patients that couldn't survive the surgery are slightly older, on average, than those who did.
14. Similarly, the Operation Year of those patients who survived the surgery is slightly greater, on average, than those who didn't.
15. The pairplot uses the three features to plot the graphs with combinations of features
16. Simply by looking, we can see that the plot between 'Operation\_Year' and 'Axillary\_Nodes' does the best job at classification
17. Although the classification is not perfect, it is better than the rest of the plots.

